

Predicting Fuel Efficiency in Car Models

The main goal of the project was to predict fuel efficiency of car models based on historical data from 2015 -2017 collected by the department of energy that can be found at

<https://www.fueleconomy.gov/feg/download.shtml>

Results showed that the best model was generated with gradient boosting using decision trees where Cross Validation score was 0.987.

Part 1

Set up:

Training and test sets for the assignment were developed based on the csv files of Fuel Economy Guide.

Training set consists of data from 2015, 2016, 2017 years, while the test set covers 2018 year data.

Predict variable: "Comb Unrd Adj FE - Conventional Fuel"

Preprocessing steps:

- deleting any columns and rows with 100% missing data.
- deleting any columns with more than 60% of missing data.
- Getting rid of the features that are correlated with the predict variable. Any features containing ['FE', 'MPG', 'EPA', 'CO2', 'Guzzler', 'cost ', 'range', 'GHG', 'Smog'] in its name were deleted.*
- Deleting all the description variables (containing 'Desc' in their names) since they don't add any value and are irrelevant for the purposes of the analysis.
- identifying remaining features types as categorical, boolean or numerical.
- converting 'Release Date' into unix timestamp.

Dealing with missing values:

To deal with missing values in the numerical data, in both training and test datasets, we computed distances between all the points, and filled missing values with the mean of the three non missing closest neighbors values.

We considered missing values in categorical as a category itself.

Scaling:

Data in both training and test sets was scaled, allowing for better comparisons and smoothing the scales.

Dealing with categorical variables:

We generated dummy variables for both the training and test sets, and filtered obtained features by selecting only the ones shared by both sets.

We divided our training set into X_{train} and y_{train} , where y_{train} contains the objective variable and X_{train} all other features. We also added a column of 1s to X_{train} to account for the constant term in the linear regression. We then split the test set in the same way.

To build our linear model, we performed a grid search cv with three types of linear regressions: Ridge, Lasso and Elastic Net.

Our best CV score is : **0.9389** using Ridge regression and alpha = **31.6**

Our R^2 test score with that model : **0.848**

Part 2

To keep the computation time low we generated our polynomial features only on numerical data. We split the data set into categorical and numerical and work with the later to perform polynomial transformation. Based on this analysis our model performance improved to CV score of **0.97**

Part 3

We used gradient boosting to improve the model. To run it we took our unscaled imputed data from the preprocessing step. Since GradientBoosting uses many decision trees in the computations there was no need to scale the data. As a result, GradientBoosting significantly improved our model yielding a CV of: **0.962**

Thus, we decided to see if we can improve even further by using polynomial features in Gradient Boosting. The resulting CV was: **0.964**

Part 4

Based on previous tasks our best model is gradient boosting using polynomial features since it has the highest test score. Using the `get_score` method on Gradient boosting we were able to extract the most important features by weight. We then fit the Gradient Boosting again only with the relevant features to achieve a test score of 0.987 (Which is the highest score so far)

```
"final_model.booster().get_score(importance_type='weight')"
```

We then collected only the features that has weight > 5. By running Gradient Boosting only on these features we achieved a better test score of 0.9871. Further increase in the value of weight leads to decrease in test score. Which means that we found our 'sweet spot'.

*Variables that were left for analysis after deleting all the extra columns:

```
categorical = ['Mfr Name',  
              'Division',  
              'Carline',  
              'Carline Class',  
              'Verify Mfr Cd',  
              'Air Aspir Method',  
              'Index (Model Type Index)',  
              'Transmission',  
              'Trans',  
              'Drive Sys',  
              'Fuel Usage - Conventional Fuel',  
              'Fuel Unit - Conventional Fuel',  
              'Fuel Metering Sys Cd',  
              'Car/Truck Category - Cash for Clunkers Bill.',  
              'Oil Viscosity']
```

```
numerical = ['Model Year',  
            'Eng Displ',  
            '# Cyl',  
            'Intake Valves Per Cyl',
```

```
'# Gears',  
'Annual Fuel1 Cost - Conventional Fuel',  
'Exhaust Valves Per Cyl',  
'$ You Spend over 5 years (increased amount spent in fuel costs over 5 years - on  
label) ',  
'Max Ethanol % - Gasoline',  
'Release Date']
```

```
boolean = ['Lockup Torque Converter',  
'Trans Creeper Gear',  
'Cyl Deact?',  
'Var Valve Timing?',  
'Var Valve Lift?',  
'Camless Valvetrain (Y or N)',  
'Stop/Start System (Engine Management System) Code',  
'Suppressed?',  
'Police/Emerg?',  
'Label Recalc?',  
'Unique Label?']
```

```
predict = ["Comb Unrd Adj FE - Conventional Fuel"]
```